

Contents

1	Fundamentals	2
1.1	Basic Statistics	2
1.2	Data Manipulation	2
1.3	Data Visualization	3
1.4	Statistical Learning Methods	4
2	Linear Regression	4
2.1	Simple Linear Regression	4
2.1.1	Formulae	5
2.1.2	Hypothesis Testing	5
2.1.3	Confidence Interval	5
2.1.4	Model Diagnostics	6
2.2	Multiple Linear Regression	7
3	<i>K</i>-means Clustering	8
3.1	Algorithm	8
3.2	Determining the Number of Clusters	9
3.3	Feature Selection	10
4	<i>K</i>-nearest Neighbor Classification	10
4.1	Algorithm	10
4.2	Examples	11
4.3	Diagnostics of Classifiers	11
4.4	<i>N</i> -Fold Cross-Validation	12
4.4.1	Algorithm	12
5	Decision Trees	15
5.1	Introduction	15
5.2	Algorithm: Entropy	16
5.3	Algorithm: Gini Index	17
6	Naïve Bayes	17
6.1	Introduction	17
6.2	Theory	18
7	Apriori Algorithm	18
7.1	Introduction	18
7.2	Theory	21
8	Additional Resources	21
8.1	Hypothesis Testing	21

1 Fundamentals

1.1 Basic Statistics

$$\text{mean}(x) = \text{sum}(x)/\text{length}(x) = \frac{1}{N} \sum_{i=1}^N x_i$$

$$\text{median}(x) = \begin{cases} X_{(\frac{N+1}{2})} & \text{if } N \text{ is odd} \\ \frac{X_{(\frac{N}{2})} + X_{(\frac{N}{2}+1)}}{2} & \text{if } N \text{ is even} \end{cases}$$

$$\text{var}(x) = \frac{1}{N-1} \sum_{i=1}^N (x_i - \bar{x})^2$$

$$\text{sd}(x) = \sqrt{\text{var}(x)}$$

$$\text{range}(x) = \max(x) - \min(x)$$

$$\text{cov}(x, y) = \frac{1}{N-1} \sum_{i=1}^N (x_i - \bar{x})(y_i - \bar{y})$$

$$\text{cor}(x, y) = \frac{\text{cov}(x, y)}{\text{sd}(x)\text{sd}(y)}$$

*Note that var, cov, and cor are calculated from sample, not population.

1.2 Data Manipulation

```

1 # import a CSV file of the total annual sales
2 setwd("c:/data")
3 sales <- read.csv("yearly_sales.csv")
4
5 # extract data
6 sales$gender      # extract gender data
7 sales[,4]         # extract 4th dataframe column
8 sales[1:2,]       # extract 1st two rows of dataframe
9 sales[,c(1,3,4)]  # extract 1st, 3rd and 4th columns of dataframe
10 # extract total sales and gender columns
11 sales[,c("sales_total", "gender")]
12 # extract all records whose gender is female
13 sales[sales$gender=="F",]
14 # extract all records with total sales above 500
15 sales[sales$sales_total>500,]
16
17 # display all the different categories in a variable
18 levels(sales$gender)
19 levels(sales[,4])

```

1.3 Data Visualization




Figure 1: Scatter plot

```

1 # simple scatter plot of number of sales versus total sales
2 png("scatter_sales.png")
3 plot(x=sales$sales_total, y=sales$num_of_orders, xlab="sales total",
4       ylab="number of orders", col="red", pch=1, main="Title")
5       # lab = label, col = color,
6       # pch = point symbols, main = title
7 dev.off() # shuts down devices

```




Figure 2: Histogram

```

1 # histogram of total sales
2 png("hist_sales.png")
3 hist(x=sales$sales_total, breaks=100, col="red")
4       # breaks = bins, more breaks more boxes
5 dev.off()

```

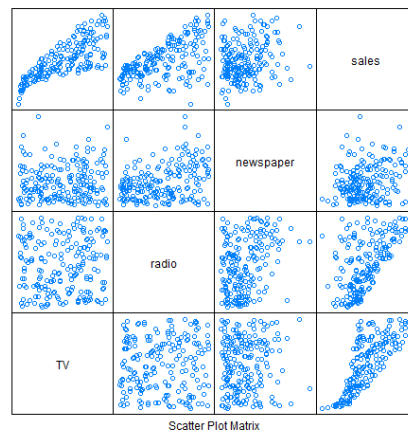


Figure 3: Pair-wise scatter plots

```

1 advert = read.csv("Advertising.csv")
2 png("scatter.png")
3 # load package before using the splom function
4 library(lattice)
5 # skip ID line, so just columns 2,3,4,5 in data
6 splom(~advert[,c(2:5)], groups=NULL, data=advert, axis.line.tck=0,
7       axis.text.alpha=0)
8 dev.off()

```

1.4 Statistical Learning Methods

Supervised	Unsupervised
Linear regression	k -means
Decision trees	Association rules
k -nearest neighbor	Hierarchical clustering
Linear discriminant analysis	Deep belief nets
Naive Bayes	Self-organizing maps

- Supervised learning: making predictions about the outcome y based on a number of predictors
- Unsupervised learning: inferring hidden structure based on data without the outcome y ('un-labeled' data)

2 Linear Regression

2.1 Simple Linear Regression

```

1 # basic summary of beta values
2 lm(sales~TV, data=advert)
3

```

```

4 # create a summary of t-statistics
5 linear.reg <- lm(sales~TV, data=advert)
6 summary(linear.reg)
7
8 # more precise values of t-statistics
9 > summary(linear.reg)$coef
10
11           Estimate Std. Error  t value    Pr(>|t|)
12 (Intercept)  7.03259355  0.457842940  15.36028 1.40630e-35
13 TV           0.04753664  0.002690607  17.66763 1.46739e-42

```

2.1.1 Formulae

$$\hat{y} = \hat{\beta}_0 + \hat{\beta}_1 x + \epsilon, \text{ where } \epsilon \sim \mathcal{N}(0, \sigma^2)$$

$$\text{Residual Sum of Squares (RSS)} = \sum_{i=1}^n [y_i - (\hat{\beta}_0 + \hat{\beta}_1 x_i)]^2$$

$$\hat{\beta}_1 = \frac{\sum_{i=1}^n x_i y_i - \bar{y} \sum_{i=1}^n x_i}{\sum_{i=1}^n x_i^2 - \bar{x} \sum_{i=1}^n x_i}$$

$$\hat{\beta}_0 = \bar{y} - \hat{\beta}_1 \bar{x}$$

* Std. Error = Standard Error of the Mean (SE).

** The z -distribution, also called the standard normal distribution, is a special normal distribution where $\mu = 0$ and $\sigma = 1$.

*** See https://en.wikipedia.org/wiki/Student's_t-distribution for the definition of t -distribution and <https://stattrek.com/probability-distributions/t-distribution.aspx> for more info.

2.1.2 Hypothesis Testing

We test the null hypothesis of $H_0 : \beta_1 = 0$ versus $H_a : \beta_1 \neq 0$ and compute the t -statistic given by:

$$t = \frac{\hat{\beta}_1 - 0}{\text{SE}(\hat{\beta}_1)}$$

where the statistic t has a t -distribution with $(n-2)$ degrees of freedom for simple linear regression. Then we can compute the p -values directly from the t -statistic or use the `pt()` function in R.

- If $p\text{-value} < \alpha$, then we reject H_0 and conclude that there is a linear relationship between x and y at the α significance level.

2.1.3 Confidence Interval

```

1 > linear.reg <- lm(sales~TV, data=advert)
2 > new_pt = data.frame(TV=200)
3 > predict(linear.reg, new_pt, level=0.95, interval="confidence")
4           fit          lwr          upr
5 1 16.53992 16.00567 17.07418

```

```

6
7 > predict(linear.reg, new_pt, level=0.95, interval="prediction")
8       fit      lwr      upr
9 1 16.53992 10.09162 22.98822

```

2.1.4 Model Diagnostics

- The quality of a linear regression fit can be assessed using the Residual Standard Error (RSE).

$$\text{RSE} = \sqrt{\frac{1}{n-2} \text{RSS}}$$

- Total Sum of Squares (TSS) measures the total variance in the response Y .

$$\text{TSS} = \sum_{i=1}^n (y_i - \bar{y})^2$$

- The R^2 statistic is another measure of the fit of the model. Larger R^2 indicates better model fit. Note that $0 \leq R^2 \leq 1$.

$$R^2 = \frac{\text{TSS} - \text{RSS}}{\text{TSS}} = 1 - \frac{\text{RSS}}{\text{TSS}}$$

```

1 # from summary(linear.reg),
2 ...
3 Residual standard error: 3.259 on 198 degrees of freedom
4 Multiple R-squared:  0.6119,    Adjusted R-squared:  0.6099
5 F-statistic: 312.1 on 1 and 198 DF,  p-value: < 2.2e-16

```

We have assumed that the error term ϵ is normally distributed with $\mu = 0$ and constant variance σ^2 . There are three ways to check:

- Plot residuals against fitted values, see if the residuals are observed somewhat evenly on both sides of the reference zero line, and the spread of the residuals is fairly constant from one fitted value to the next.




Figure 4: Residuals against fitted values

```

1 > png("residuals.png")
2 > plot(x=linear.reg$fitted.values, y=linear.reg$residuals, xlab=
  "Fitted values", ylab="Residuals", col="red")
3 > abline(0,0) # sketch the line y=ax+b: abline(a,b)
4 > dev.off()

```

2. Histogram of residuals to check the normality assumption for the error terms.




Figure 5: Histogram of residuals

```

1 > png("hist_residuals.png")
2 > hist(x=linear.reg$residuals, breaks=50, col="red")
3 > dev.off()

```

3. Quantile-Quantile (QQ) plot that compares the observed residuals against the quantiles of the theoretical normal distribution.

```

1 > png("qq_residuals.png")
2 > qqnorm(linear.reg$residuals, ylab="Residuals", main=" ")
3 > qqline(linear.reg$residuals)
4 > dev.off()

```

2.2 Multiple Linear Regression

```

1 > resale <- read.csv("hdbresale_reg.csv")
2 > out <- lm(resale_price~town+floor_area_sqm+flat_type, data=resale)
3
4 # compute confidence intervals on the parameters
5 > confint(out, level=0.95)
6
7 # provide confidence interval on expected predicted values

```

```

8 > new_pt <- data.frame(town="CENTRAL AREA", flat_type="3 ROOM",
   floor_area_sqm=70)
9 > predict(out, new_pt, level=0.95, interval="confidence")
10 > predict(out, new_pt, level=0.95, interval="prediction")

```

In a model with p input variables,

$$y = \beta_0 + \beta_1 x(1) + \beta_2 x(2) + \cdots + \beta_p x(p) + \epsilon,$$

where $x(j)$ are the input variables, $j = 1, 2, \dots, p$

3 K -means Clustering




Figure 6: K -means clustering

```

1 resale <- read.csv("hdbresale_cluster.csv")
2 kout <- kmeans(resale[,c("floor_area_sqm", "amenities")], centers=2)
3 plot(resale$floor_area_sqm, resale$amenities,
4       col=kout$cluster,
5       xlab="Floor area in sqm",
6       ylab="Number of amenities", pch=19)

```

3.1 Algorithm

1. Choose the value of k and the k initial guesses for the centroids
2. Compute the Euclidean distance from each data point to each centroid
3. Assign the points to their closest centroids
4. Compute the centroids in each cluster
5. Repeat steps 2-4 until the algorithm converges to an answer

3.2 Determining the Number of Clusters

For a given point z_i at $(x(1)_i, x(2)_i, \dots, x(p)_i)$ and a centroid D at $(x(1)_d, x(2)_d, \dots, x(p)_d)$, the distance between z_i and D is

$$\text{dist}(z_i, D) = \sqrt{\sum_{j=1}^p (x(j)_i - x(j)_d)^2}$$

The centroid for a cluster of m points, $(x(1)_i, x(2)_i, \dots, x(p)_i)$ for $i = 1, 2, \dots, m$ is given by

$$\left(\frac{1}{m} \sum_{i=1}^m x(1)_i, \frac{1}{m} \sum_{i=1}^m x(2)_i, \dots, \frac{1}{m} \sum_{i=1}^m x(p)_i \right)$$

For M data points z_1, z_2, \dots, z_M with p features the Within Sum of Squares (WSS) is calculated via

$$\begin{aligned} \text{WSS} &= \sum_{i=1}^M \text{dist}(z_i, D_i)^2 \\ &= \sum_{i=1}^M \sum_{j=1}^p (x(j)_i - x(j)_{D_i})^2 \end{aligned}$$

```
1 > grade_input = read.csv("grades_kmean_input.csv")
2 > kout <- kmeans(grade_input[,c("English", "Math", "Science")],
3   centers = 3)
4 > kout$withinss
[1] 34806.339 22984.131 6692.589
```




Figure 7: WSS against the number of clusters (Elbow curve)

```

1 wss <- numeric(15) # creates an array of size 15
2 for (k in 1:15) {
3   km <- kmeans(grade_input[,c("English", "Math", "Science")],
4     centers = k, nstart=25)
5   # nstart = do 25 times of K-means clustering
6   wss[k] <- sum(km$withinss) # km$tot.withinss
7 }
8 plot(1:15, wss, type="b", xlab="Number of Clusters", ylab="Within
9   Sum of Squares")

```

3.3 Feature Selection

- Include only important features, too many attributes can minimize the impact of the most important variables
- Identify any highly correlated attributes and only use one or two of them
- Choose the appropriate unit of measure for each attribute
- Rescale some features so that one feature does not have disproportionate effects on results

4 K -nearest Neighbor Classification

4.1 Algorithm

1. For two given data points in p -dimensional feature space, z_i at $(x(1)_i, x(2)_i, \dots, x(p)_i)$ and z_j at $(x(1)_j, x(2)_j, \dots, x(p)_j)$, the Euclidean distance between z_i and z_j is

$$\text{dist}(z_i, z_j) = \sum_{l=1}^p (x(l)_i - x(l)_j)^2$$

Find the k nearest neighbors to the data point in the feature space in terms of Euclidean distance

2. The fitted \hat{Y} with k nearest neighbors for a new data point with feature values x is

$$\hat{Y} = \frac{1}{k} \sum_{z_i \in N_k(x^*)} y_i$$

The prediction error of the model can be decomposed into

$$\text{error} = \text{bias}^2 + \text{variance} + \text{irreducible error}$$

When k increases, the variance decreases but bias increases, this is known as the bias-variance tradeoff. There is an inverse relationship between variance and biased

4.2 Examples

- Example 1: The Stock Market Data

```

1 > market <- read.csv("Smarket.csv")
2
3 > train = (market$Year < 2005)
4 > train.data = market[train,]
5 > test.data = market[!train,]
6
7 > train.x = train.data[,c("Lag1", "Lag2", "Lag3", "Lag4", "Lag5")]
8 > test.x = test.data[,c("Lag1", "Lag2", "Lag3", "Lag4", "Lag5")]
9 > train.y = train.data[,c("Direction")]
10 > test.y = test.data[,c("Direction")]
11
12 > library(class) # knn() is a function in 'class'
13 # knn() needs at least 4 args
14 > knn.pred = knn(train.x, test.x, train.y, k=1)
15 # create a table for diagnostic of classifier
16 > confusion.matrix = table(knn.pred, test.y)
17 > confusion.matrix
18      test.y
19 knn.pred Down Up
20      Down   55 66
21      Up    56 75

```

- Example 2: Caravan Insurance Data

```

1 > caravan = read.csv("Caravan.csv")
2 # exclude ID column, Purchase (Y/N) column
3 > caravan = caravan[, -1]
4 > standardized.X = scale(caravan[, -86])
5 # every column of standardized.X has var=1 and mean=0

```

Precision is used as a diagnostic tool here because campaign can be more targeted to potential customers who are more likely to purchase insurance

- Example 3: Customer Churn

Precision as a metric is most useful here because companies can target customers that are more likely to churn by offering various incentives

4.3 Diagnostics of Classifiers

- The *accuracy* (overall success rate) is a metric defining the rate at which a model has classified the records correctly

$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN} \times 100\%$$

- The true positive rate (TPR) shows the proportion of positive instances the classifier correctly identified

$$\text{TPR} = \frac{TP}{TP + FN}$$

- The false positive rate (FPR) shows what percentage of negatives the classifier marked as positive. Also called the false alarm rate or the type I error rate

$$\text{FPR} = \frac{FP}{FP + TN}$$

- The false negative rate (FNR) shows what percentage of positives the classifier marked as negative. Also called the miss rate or the type II error rate

$$\text{FNR} = \frac{FN}{TP + FN}$$

- *Precision* is the percentage of instances marked positive that really are positive

$$\text{Precision} = \frac{TP}{TP + FP}$$

4.4 N-Fold Cross-Validation

4.4.1 Algorithm

1. The entire dataset is randomly split into N datasets of approximately equal size.
2. $N - 1$ of these datasets are treated as the training dataset, while the remaining one is the test dataset. A measure of the model error is obtained.
3. This process is repeated across the various combinations of N datasets taken $N - 1$ at a time.
4. The observed N models errors are averaged across the N folds.

```

1 > iris = read.csv("iris.csv", header=FALSE)
2 > names(iris) = c("X1", "X2", "X3", "X4", "Y")
3 > library(class)
4 > X = iris[,c("X1", "X2", "X3", "X4")]
5 > Y = iris[,c("Y")]
6
7 # 150 datapoints are randomly divided into 10 sets
8 > n_folds=10
9 > folds_j <- sample(rep(1:n_folds, length.out = 150))
10 > table(folds_j)
11 folds_j
12  1  2  3  4  5  6  7  8  9 10
13 15 15 15 15 15 15 15 15 15 15
14 > test_j <- which(folds_j == 1)
15
16 > pred <- knn(train=X[-test_j,], test=X[test_j,], cl=Y[-test_j], k
    =1)

```

```

17 > mean(pred == Y[test_j]) # proportion of datapoints predicted
    correctly
18 [1] 0.9333333
19 > mean(pred != Y[test_j]) # proportion of datapoints predicted
    incorrectly
20 [1] 0.06666667
21 }

```




Figure 8: KNN: Error rate against different values of K

```

1 n_folds=10
2 folds_j <- sample(rep(1:n_folds, length.out = 150))
3
4 error_differentK = numeric(50)
5 accur_differentK = numeric(50)
6
7 for (K in 1:50) {
8
9   error = numeric(10)
10  accur = numeric(10)
11  for (j in 1:10) {
12    test_j <- which(folds_j == j)
13    pred <- knn(train=X[-test_j,], test=X[test_j,], cl=Y[-test_j], k=K
14    )
15    error[j] = mean(pred != Y[test_j])
16    accur[j] = mean(pred == Y[test_j])
17  }
18  error_differentK[K] = mean(error)

```

```

18 accur_differentK[K] = mean(accur)
19
20 plot(1:50, error_differentK, type="o", ylab="error rate", xlab="K",
      cex.axis=1, cex=2)
21 plot(1:50, accur_differentK, type="o", ylab="accuracy", xlab="K",
      cex.axis=1, cex=2)

```

Using the caret library:




Figure 9: *N*-fold cross validation using caret

```

1 library(caret)
2
3 fitControl <- trainControl(## 10-fold cross validation
4                             method="repeatedcv",
5                             number=10,
6                             ## repeated 10 times
7                             repeats=10)
8 set.seed(2)
9 knnFit1 <- train(Y ~., # predict Y from everything else
10                 data=iris, method="knn",
11                 trControl=fitControl, #trainControl
12                 preProcess=c("center","scale"), # centered & scaled
13                 tuneLength=50)
14 plot(knnFit1)

```

5 Decision Trees

5.1 Introduction

'Branch' refers to the outcome of a decision and is visualized as a line connecting two nodes. The branching of a node is referred to as a split.

'Internal nodes' are the decision or test points. The topmost internal node is called the root.

The **'depth'** of a node is the minimum number of steps required to reach the node from the root.

Leaf nodes are at the end of the last branches on the tree

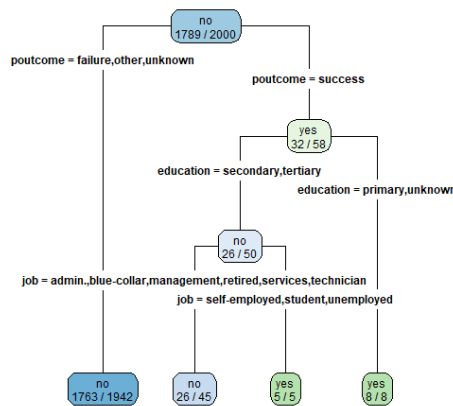


Figure 10: Decision Tree

```

1 install.packages("rpart")
2 install.packages("rpart.plot")
3 # rpart contains functions for modelling decision trees
4 # rpart.plot enables the plotting of a tree
5 library("rpart")
6 library("rpart.plot")
7
8 bankdata = read.csv("bank-sample.csv")
9
10 fit <- rpart(subscribed~job + marital + education +
11             default + housing + loan + contact + poutcome,
12             method="class", data=bankdata,
13             control=rpart.control(minsplit=1, maxdepth=4),
14             # minsplit =
15             # maxdepth = max depth of tree
16             parms=list(split="information"))
17             # "information" => entropy algo
18             # "gini"       => gini algo
19 rpart.plot(fit, type=4, extra=2, clip.right.labs=FALSE,
20           varlen=0, faclen=0)

```

5.2 Algorithm: Entropy

- The **purity** of a node is defined as its probability of the corresponding class
For example, in the root node, $P(\text{subscribed} = 0) = \frac{1789}{2000}$. Hence the root is 89.45% pure on the subscribed = 0 class and 10.55% pure on the subscribed = 1 class
- Given variable Y and the set of possible categorical values it can take, (y_1, y_2, \dots, y_K) , the **entropy** of Y is defined as

$$D_Y = - \sum_{j=1}^K P(Y = y_j) [\log_2 P(Y = y_j)]$$

where $P(Y = y_j)$ denotes the purity or the probability of the class $Y = y_j$ and $\sum_{j=1}^K P(Y = y_j) = 1$.

- Heuristically, entropy is a measure of unpredictability. Lower entropy \Rightarrow lower uncertainty and likewise, higher entropy \Rightarrow higher uncertainty
- The base entropy is defined as the entropy of the output variable at the root node
- Suppose we have a feature variable X and the split values (x_1, x_2) . The **conditional entropy** given feature X and the split points (x_1, x_2) is defined as

$$\begin{aligned} D_{Y|X} &= \sum_{i=1}^2 P(X = x_i) \cdot D(Y|X = x_i) \\ &= - \sum_{i=1}^2 \left\{ P(X = x_i) \sum_{j=1}^K P(Y = y_j | X = x_i) [\log_2 P(Y = y_j | X = x_i)] \right\} \end{aligned}$$

- The reduction in entropy from the base entropy is also known as **information gain**.

Entropy algorithm (from 8.1 slide 38):

1. Therefore, the decision tree algorithm proceeds at the root node by calculating the conditional entropy for (i) each feature variable X and (ii) its different split points
2. Then, the decision variable and its split points are selected based on the largest information gain (or decrease from base entropy)
3. At internal nodes, the decision tree algorithm proceeds similarly by calculating the conditional entropy for (i) each feature variable X and (ii) its different split points.
4. However, the sample for calculating the base and conditional entropies is restricted to the one at the node.
5. The tree is built recursively until a criteria is met, for example:
 - (a) All the leaf nodes in the tree satisfy the minimum purity threshold.
 - (b) The tree cannot be further split with the preset minimum purity threshold.
 - (c) Any other stopping criterion is satisfied (such as the maximum depth of the tree).

5.3 Algorithm: Gini Index

- Given variable Y and the set of possible categorical values it can take, (y_1, y_2, \dots, y_K) , the Gini index of Y is defined as

$$G_Y = \sum_{j=1}^K P(Y = y_j)[1 - P(Y = y_j)]$$

where $P(Y = y_j)$ denotes the purity or the probability of the class $Y = y_j$ and $\sum_{j=1}^K P(Y = y_j) = 1$.

6 Naïve Bayes

6.1 Introduction

```

1 setwd("C:/Users/Claudeon/Downloads")
2 fruit.dat <- read.csv("fruit.csv")
3 fruit.dat <- fruit.dat[, -1]
4 fruit.dat <- data.frame(lapply(fruit.dat, as.factor))
5                               # change df to factor
6 # import library!
7 library(e1071)
8
9 model <- naiveBayes(Fruit~Long+Yellow+Sweet, fruit.dat)
10
11 newdata <- data.frame(Long=1, Sweet=1, Yellow=0)
12 newdata <- data.frame(lapply(newdata, as.factor))
13
14 > predict(model, newdata, "raw")
15      Banana      Orange      Other
16 [1,] 0.5284404 0.0587156 0.412844
17
18 > predict(model, newdata, "class")
19 [1] Banana
20 Levels: Banana Orange Other
21
22 # Other Shortcuts!!
23 > test <- table(fruit.dat[, c("Fruit", "Long")])
24 > test
25      Long
26 Fruit    0    1
27 Banana 300 200
28 Orange 280  20
29 Other  100 100
30 > rowSums(test)
31 Banana Orange  Other
32   500    300    200
33 > test/rowSums(test)

```

```

34      Long
35 Fruit      0      1
36 Banana 0.60000000 0.40000000
37 Orange 0.93333333 0.06666667
38 Other  0.50000000 0.50000000

```

6.2 Theory

- Assuming conditional independence,

$$P(AB|C) = P(A|C)P(B|C) \Leftrightarrow P(A|BC) = P(A|C) \leftarrow \text{definition}$$

* Event B does not affect A given C

- Assigns a classified label to an object with m feature variables $X = \{X_1, X_2, \dots, X_m\}$ such that the predicted label corresponds to the largest value of $P(Y = y_j|X)$, $j = 1, 2, \dots, k$ and $Y \in \{y_1, y_2, \dots, y_k\}$

$$\begin{aligned}
 P(Y = y_j|X) &= \frac{P(X_1 = x_1, X_2 = x_2, \dots, X_m = x_m|Y = y_j) \times P(Y = y_j)}{P(X_1 = x_1, X_2 = x_2, \dots, X_m = x_m)} \\
 &= \frac{P(Y = y_j) \times \prod_{i=1}^m P(X_i = x_i|Y = y_j)}{P(X_1 = x_1, X_2 = x_2, \dots, X_m = x_m)} \\
 &\propto P(Y = y_j) \times \prod_{i=1}^m P(X_i = x_i|Y = y_j)
 \end{aligned}$$

$$\Leftrightarrow \log P(Y = y_j|X) = \log P(Y = y_j) + \sum_{i=1}^m \log P(X_i = x_i|Y = y_j)$$

- Ignore denominator term since probability is constant

7 Apriori Algorithm

7.1 Introduction

```

1 library("arules")
2 library("arulesViz")
3
4 data(Groceries) # import Groceries from arules
5                 # class ==> transaction
6 # to show/inspect, use Groceries or summary()
7 > Groceries@itemInfo[1:5,] # shows all the items and their
   categories
8   labels level2      level1
9 1 frankfurter sausage meat and sausage
10 2   sausage sausage meat and sausage
11 3  liver loaf sausage meat and sausage
12 4     ham sausage meat and sausage

```

```

13 5      meat sausage meat and sausage
14 > Groceries@data[10:15,10:15]  # shows items and which
15                                # transactions they're in
16                                # in sparse format
17 6 x 6 sparse Matrix of class "ngCMatrix"
18
19 [1,] . . . . .
20 [2,] . . . | . .
21 [3,] . . . . .
22 [4,] . . . . .
23 [5,] . . | . . .
24 [6,] . | | . . |
25
26 # convert sparse format to namelist
27 # note that column corresponds to transaction no.
28 #      row corresponds to item no.
29 > apply(Groceries@data[,101:105], 2, function(r)paste(
30   Groceries@itemInfo[r, "labels"], collapse=", "))
31 [1] "soda, misc. beverages"
32 [2] "sausage, pork, grapes, whole milk, rolls/buns, pastry, soda,
33   specialty bar, bathroom cleaner"
34 [3] "frankfurter, rolls/buns, bottled water"
35 [4] "sausage, whole milk, yogurt, coffee, fruit/vegetable juice,
36   bottled beer, softener, napkins, photo/film, shopping bags"
37 [5] "soda"
38
39 # Apriori Algorithm
40 > itemsets <- apriori(Groceries, parameter=list(minlen=1, maxlen=1,
41   support=0.02, target="frequent itemsets"))
42   # min/maxlen = min/max length of itemsets
43   # target = type of association mined
44 # inspect items
45 > inspect(head(sort(itemsets, by="support"), 5))
46
47   items          support  count
48 [1] {whole milk}    0.2555160 2513
49 [2] {other vegetables} 0.1934926 1903
50 [3] {rolls/buns}    0.1839349 1809
51 [4] {soda}          0.1743772 1715
52 [5] {yogurt}        0.1395018 1372
53
54 # recommended apriori implementation
55 > itemsets <- apriori(Groceries, parameter=list(minlen=1, support
56   =0.02, target = "frequent itemsets"))
57 > summary(itemsets)
58 set of 122 itemsets
59
60 most frequent items:
61   whole milk other vegetables          yogurt      rolls/buns

```

```

56          28          20          11          10
57          soda          (Other)
58          10          108
59
60 element (itemset/transaction) length distribution:sizes
61  1  2  3
62 59 61  2
63 ...

```




Figure 11: Association rule scatterplot

```

1 > rules <- apriori(Groceries, parameter=list(support =0.001,
2         confidence=0.6, target="rules"))
3 # lower min support allows more rules to show up
4 # the highest lift occurs at a low support and a low confidence.
5 > plot(rules)
6
7 # sort association rules by lift
8 > inspect(head(sort(rules, by="lift"), 3))
9
10 lhs          rhs          support
11 confidence
12 [1] {Instant food products, soda} => {hamburger meat} 0.001220132
13 0.6315789
14 [2] {soda, popcorn} => {salty snack} 0.001220132
15 0.6315789
16 [3] {ham, processed cheese} => {white bread} 0.001931876
17 0.6333333
18 coverage lift count
19 [1] 0.001931876 18.99565 12

```

```

14 [2] 0.001931876 16.69779 12
15 [3] 0.003050330 15.04549 19

```

7.2 Theory

- Given a k -itemset $L = \{i_1, i_2, \dots, i_k\}$, the **support** of L is the percentage of transactions that contain L .
- (Apriori/Downward Closure Property)** A **frequent** itemset has items that appear together often enough (when support > minimum support criterion)
 - Note that if $X \subseteq Y$ and Y is a frequent itemset then X is also a frequent itemset.

- Confidence** is defined as the measure of certainty or trustworthiness associated with each discovered rule.

$$\text{Confidence}(X \rightarrow Y) = \frac{\text{Support}(X \wedge Y)}{\text{Support}(X)}$$

- A relationship is thought of as interesting when the algorithm identifies the relationship with a measure of confidence greater than or equal to a predefined threshold.
- Lift** measures how many times more often X and Y occur together than expected if they are statistically independent of each other ($\text{Lift} = 1$ means X and Y are statistically independent of each other, the higher the value the more useful the rule is)

$$\text{Lift}(X \rightarrow Y) = \frac{\text{Support}(X \wedge Y)}{\text{Support}(X) \times \text{Support}(Y)}$$

- Leverage** ($\text{Leverage} = 0$ means X and Y are statistically independent of each other, larger value indicates stronger relationship)

$$\text{Leverage}(X \rightarrow Y) = \text{Support}(X \wedge Y) - \text{Support}(X) \times \text{Support}(Y)$$

8 Additional Resources

8.1 Hypothesis Testing

- Steven Walker. 2020. *A Level Maths - hypothesis tests and the art of being non-assertive*. Source: <https://www.ocr.org.uk/blog/a-level-maths-hypothesis-tests-and-the-art-of-being-non-as>
- Khan Academy. 2018. *Calculating t statistic for slope of regression line*. Source: <https://www.youtube.com/watch?v=7MAuojBTF-g>
- Khan Academy. 2018. *Using a P-value to make conclusions in a test about slope*. Source: <https://www.youtube.com/watch?v=Mpd83AuDTrU>